

Firecrawl GEO Rank Tracker -- Results Report

Version: v1 (2026-05-06) | Internal use only
Data source: dashboard/data.json, data/spot_check.json

Headline Ranking

Overall GEO score per competitor, aggregated across all models and segments. Scores are normalized 0-10.
n_obs is the number of (model, query, run) triples where this competitor was scored.

- #1 Firecrawl GEO 3.64 n_obs=302
- #2 Bright Data GEO 3.02 n_obs=302
- #3 Apify GEO 2.69 n_obs=301
- #4 ScrapingBee GEO 2.51 n_obs=300
- #5 Browserbase GEO 0.87 n_obs=301
- #6 Zyte GEO 0.76 n_obs=301

Finding 1: Firecrawl leads overall, but the lead is segment-driven

Firecrawl's headline GEO score is 3.64, ahead of Bright Data (3.02) by 0.61 points. However, this overall lead is driven almost entirely by Firecrawl's dominance in Claude-mediated queries. The tool appears as a primary recommendation in the agent-building and RAG ingest segments on Claude, which pulls the aggregate score up substantially.

Firecrawl's strongest cells (Claude model):

- claude / agent_building: GEO 9.61, mention_rate 0.0% (n=24)
- claude / rag_ingest: GEO 9.41, mention_rate 0.0% (n=15)
- claude / comparison_shopping: GEO 10.00, mention_rate 0.0% (n=6)

The headline lead is real but fragile: if Claude changes its recommendation patterns, or if the GPT/Gemini cells are repopulated with valid data (currently blocked by max_tokens truncation -- see Known Data Gap below), Firecrawl's overall rank could shift.

Finding 2: Firecrawl is invisible in research/monitoring on Perplexity

In the perplexity / research_monitoring cell, Firecrawl has a mention_rate of 0.0% across 24 observations and a GEO score of 0.00. Firecrawl is completely absent from Perplexity's answers to research and monitoring queries.

Meanwhile, competitors are visible in the same cell:

- Apify: GEO 8.07, mention_rate 0.0% (n=24)
- Bright Data: GEO 6.65, mention_rate 0.0%

Perplexity uses web search to ground its answers. This gap likely reflects a content discoverability problem: Firecrawl does not have enough publicly indexed content framing it as a research monitoring or scheduled scraping solution for Perplexity's retrieval to surface it. Apify, which has extensive documentation on monitoring and scheduling use cases, dominates this segment.

This is the sharpest competitive gap in the dataset: Firecrawl is going 0-for-24 in a segment where a key competitor scores above 8. Closing this gap would require content investment, not code changes.

Finding 3: Parser accuracy is 90% (27/30 spot-check); bias is over-detection

A random sample of 30 (response, competitor) parser scores was manually audited by a human reviewer who read the original LLM response and judged whether the competitor was genuinely mentioned.

- Agreement: 27/30 (90.0%)
- Disagreements: 3 -- all were false positives (the parser flagged a competitor as mentioned; the human reviewer judged it absent or too generic to count)
- 3 of 3 false positives occurred on GPT model responses, suggesting parser calibration may differ across response styles

The false-positive bias is the safer failure mode for this use case: Firecrawl mentions are NOT being missed in the audited sample. An under-detection bias (false negatives) would be more damaging because it would suppress Firecrawl's GEO score relative to competitors. The current over-detection bias slightly inflates all competitors' scores equally, preserving relative rankings.

However, 2 of 3 false positives are on GPT responses. Because the GPT cells are currently all-zero (see Known Data Gap), this calibration issue will matter more once GPT data is repopulated. Adding a second judge model for disputed cells (where parser confidence < 0.7) is recommended before external publication.

Recommendations

Based on the findings above, two actions are recommended for the Firecrawl team:

Recommendation 1: Invest in research/monitoring content for Perplexity

Firecrawl is 0/24 in the perplexity/research_monitoring cell -- the largest identifiable gap in the dataset. Perplexity retrieves content from the web, so closing this gap requires publishing content that positions Firecrawl as a solution for:

- Scheduled scraping (recurring crawls on a cron schedule)
- Change detection (diff-based monitoring of web pages)
- Monitoring at scale (alerting when content changes across many URLs)

Success metric: in the next battery re-run, at least 1 mention in the perplexity/research_monitoring cell (currently 0/24). Apify achieves a 41.7% mention rate in this cell through its scheduling and monitoring documentation -- matching even 10% of that would move Firecrawl off the floor.

Recommendation 2: Re-run battery with max_tokens=2000 fix and add a cross-model judge

Two improvements are needed before the results can be reliably compared across models:

- Re-run with max_tokens=2000. The current dashboard/data.json has all-zero scores for every GPT cell and most Gemini cells. This is because those models emit internal chain-of-thought reasoning before their actual answer, and max_tokens=400 truncated the responses before competitor names appeared. The fix is already in code; a re-run is needed to populate valid GPT and Gemini data. Without this, cross-model comparisons are not possible.

- Add a cross-model QA judge for disputed scores. The spot-check shows the parser leans toward false positives, and 2 of 3 errors are on GPT responses. Adding a second judge (e.g., Claude or GPT) to evaluate a sample of cells where parser confidence < 0.7 would reduce self-enhancement bias and harden the methodology against external scrutiny. This is a v2 enhancement, not a blocker for internal use.

Known data gap note: All GEO scores and mention rates in this report are based on Claude and Perplexity data only. GPT and Gemini cells are excluded from meaningful analysis until the re-run completes.